

FRA503 Deep Reinforcement Learning for Robotics

Final Report: SplendorDRL with Event Based Reward Shaping and Double DQN

Name	Student ID
Chantouch Orungrote	66340500011
Sasish Kaewsing	66340500076

1 Introduction

This project develops a Deep Reinforcement Learning (DRL) agent for *Splendor*, a deterministic, perfect-information, turn-based engine-building board game. Players collect gem tokens to purchase development cards granting permanent bonuses, winning by reaching at least 15 prestige points before the opponent.

The primary motivation is *Rinascimento* [1], which shows that score-only reward signals are too sparse for effective game-playing and that event-based signals provide richer guidance. This project extends that insight into model-free DRL, investigating whether event-based reward shaping improves training, and—crucially—whether advanced architectures (Dueling DQN, Prioritized Experience Replay, self-play curriculum) overcome pathological behaviors that naive reward shaping introduces.

The system supports two operational modes: **Mode 1** (AI vs. AI) for training and evaluation, and **Mode 2** (AI vs. Human) for interactive demonstration via a real-time Pygame visualization.

2 Problem Statements

The game presents three core challenges for DRL:

1. **Sparse rewards** — most actions yield zero immediate prestige gain.
2. **Delayed strategic payoff** — gem-collection and card-purchase decisions compound over many turns before affecting the final score.
3. **State-dependent legal action space** — the valid action set changes every step based on current gem counts, visible cards, and reserved cards.

A fourth challenge—not anticipated in the proposal—emerged during training: a **discount-exploitation pathology** (Section 10) in which the agent learns to extend episode length rather than attempt to win, exploiting the γ -discounting of terminal penalties.

3 Objectives

- Adapt an open-source *Splendor* simulator into a Gym-compatible environment with legal action masking and a modular, swappable opponent interface.
- Implement a Double Deep Q-Network (DDQN) agent with optional Dueling architecture, Prioritized Experience Replay (PER), soft target updates, and masked ε -greedy exploration.
- Design and compare multiple event-based reward shaping schemes (C2–C6) against a score-only baseline (C1) and a lower bound (C0 vs. random).
- Implement a self-play curriculum to prevent overfitting to a fixed opponent strategy.
- Develop a real-time Pygame visualization supporting both AI vs. AI and AI vs. Human play modes.
- Evaluate win rate, training stability, and sample efficiency across all seven experimental conditions.

4 Scope

- Standard two-player *Splendor*: gem collection, card purchase across three tiers, card reservation, noble attraction, and prestige-based victory.
- Seven experimental conditions (C0–C6) covering reward shaping variants, opponent types, and architectural ablations.
- DDQN-based learning with discrete action formulation and legal action masking; optional Dueling head and PER.
- Full Pygame visualization system supporting both operational modes.
- Empirical comparison of all conditions over 300 000 training episodes.

5 System Overview

The system consists of five modules, all active in both modes; only the Opponent Module switches behavior.




Figure 1: System architecture. Main training path: environment → reward module → DDQN agent → replay buffer. Visualization and opponent modules support both modes; in C3 the opponent pool stores past agent snapshots.

5.1 State Representation

The observation is a fixed-length normalized vector ($\in [0, 1]^{213}$):

Group	Content	Dim.
Visible cards	12 cards $\times$ (5 gem costs + 1 prestige + 5 bonus one-hot)	132
Noble tiles	3 nobles $\times$ (5 gem requirements + 1 prestige)	18
Board supply	6 gem-color counts (including gold/wild)	6
Agent state	6 gem counts + 5 bonuses + 3 $\times$ 11 reserved cards + 1 prestige	45
Opponent state	6 gem counts + 5 bonuses + 1 prestige (public info only)	12
Total		213

5.2 Action Representation

The discrete action space (size ≤ 48) covers: gem collection (≤ 15 combinations), card purchase (up to 15 visible or reserved cards), card reservation (12 visible cards), and gem discard (6 colors, triggered when the 10-gem limit is exceeded). A binary legal mask $m \in \{0, 1\}^{48}$ is recomputed at every step.

5.3 Opponent Module

The opponent exposes a single interface `select_action(obs, mask, env)`, making it fully swappable:

- **Greedy (C0–C2, C4–C6)**: buys highest-prestige affordable card; on ties, collects gems toward cheapest visible card.
- **Human (Mode 2)**: presents legal options via the visualization; relays click as integer action index.
- **Self-Play (C3)**: randomly selects one past agent snapshot from a rolling pool of 10; falls back to greedy when pool is empty.

5.4 Pygame Visualization

The Visualization Module renders the complete game state at every turn. Figure 2 shows Mode 2 (AI vs. Human), the primary interactive mode for inspecting learned policy quality.




Figure 2: Mode 2 visualization (AI vs. Human). The board shows gem supply, three card tiers, noble tiles, and both players' states. Legal actions for the human player are highlighted in green; cards are clickable and gem combinations appear as on-screen buttons. The agent's chosen action is displayed in the status bar before application.

6 Techniques

6.1 Double DQN with Dueling Architecture

DDQN [3] decouples action selection (online network θ) from evaluation (target network θ^-):

$$y_t = r_t + \gamma Q\left(s_{t+1}, \arg \max_{a'} Q_{\text{masked}}(s_{t+1}, a'; \theta); \theta^-\right)$$

Dueling DQN [4] decomposes $Q(s, a) = V(s) + A(s, a) - \bar{A}(s)$, allowing the agent to learn state value independently of action advantage—critical in Splendor, where the board state often determines turn quality regardless of the specific action chosen.

The network architecture (with LayerNorm after each hidden layer for gradient stability):

$$\text{Input}(213) \rightarrow \text{FC}(256) + \text{LN} + \text{ReLU} \rightarrow \text{FC}(256) + \text{LN} + \text{ReLU} \rightarrow \begin{cases} \text{FC}(128) \rightarrow V(s) \\ \text{FC}(128) \rightarrow A(s, a) \end{cases}$$

Training uses Huber loss (smooth ℓ_1), gradient clipping at 1.0, and *soft* target updates: $\theta^- \leftarrow \tau \theta + (1 - \tau)\theta^-$, $\tau = 0.005$.

6.2 Experience Replay and PER

A circular buffer [2] (capacity 10^5 , min fill 5,000) stores transitions $(s_t, a_t, r_t, s_{t+1}, m_{t+1}, \text{done}_t)$. Mini-batches of 64 are sampled uniformly in baseline conditions.

Prioritized Experience Replay (PER) [5] (C3 only): transition i is sampled with probability $p_i \propto |\delta_i|^\alpha$, $\alpha = 0.6$. A SumTree provides $O(\log n)$ priority updates. Importance-sampling weights with β annealed from 0.4 to 1.0 over 500 000 steps correct the resulting bias.

6.3 Legal Action Masking

Illegal actions are excluded from both selection and bootstrapping [7]:

$$Q_{\text{masked}}(s, a; \theta) = \begin{cases} Q(s, a; \theta), & m_a = 1 \\ -10^4, & m_a = 0 \end{cases}$$

The clamp at -10^4 (not $-\infty$) prevents float16 overflow under AMP on CUDA.

6.4 ε -Greedy Exploration

Exploration is restricted to legal actions, with per-episode decay:

$$\varepsilon_{k+1} = \max(\varepsilon_{\min}, \varepsilon_k \cdot \lambda), \quad \varepsilon_0 = 1.0, \varepsilon_{\min} = 0.05, \lambda = 0.99997$$

The learning rate is stepped down by $10\times$ at episodes 200 000 and 250 000.

6.5 Event-Based Reward Shaping

Motivated by [1], the shaped reward densifies the learning signal:

$$r_t = w_{\text{gem}} \cdot \phi_{\text{gem}} + \sum_{k=1}^3 w_{L_k} \cdot \phi_{L_k} + w_P \cdot \Delta P_t + w_N \cdot \phi_N + w_{\pm} \cdot \mathbf{1}[\text{win/lose}]$$

Win bonus = +10, loss penalty = -10 for all conditions. Four profiles were implemented (Table 1).

Table 1: Reward weight profiles. C1 uses score-only ($w_{\text{gem}} = w_{L_k} = w_N = 0$, $w_P = 1$).

Cond.	Profile	w_{gem}	w_{L_1}	w_{L_2}	w_{L_3}	w_P	w_N
C2	Event-shaped	0.002	0.05	0.10	0.20	1.0	3.0
C4	Card Rush	0.000	0.15	0.25	0.40	1.0	3.0
C5	Noble Hunter	0.001	0.03	0.08	0.30	1.0	6.0
C6	Balanced Dense	0.010	0.10	0.20	0.40	1.0	4.0

6.6 Self-Play Curriculum (C3)

Every 2000 episodes the current online network is pushed into a FIFO pool (max 10 snapshots). At each episode start, one snapshot is randomly selected as the opponent. Early episodes face the greedy fallback; later episodes face progressively stronger past-self versions. This prevents the agent from learning fixed-opponent exploits (Section 10).

7 Framework

The environment is an episodic MDP $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ [6], $\gamma = 0.99$, built on a Gym-compatible interface exposing `reset()`, `step(action)`, mask generation, terminal detection, and reward computation. The opponent’s turn is executed inside `step()` so the agent observes only its own decision points.

Training (Mode 1): Initialize environment and agent with random weights; for each of 300 000 episodes: encode (s_t, m_t) ; select a_t via masked ε -greedy; receive $(s_{t+1}, r_t, \text{done})$; store transition; sample mini-batch and update θ ; soft-copy to θ^- ; decay ε at episode end. Push snapshot to self-play pool every 2000 episodes (C3 only).

Deployment (Mode 2): Load frozen weights ($\varepsilon = 0$); visualization renders the full board each turn; agent selects greedy legal action; human selects from highlighted legal options via click; game ends when either player reaches 15 prestige points.

8 Experimental Setup

Cond.	Description	Reward	Opponent	Architecture
C0	Lower bound	Score-only	Random	DDQN
C1	Baseline	Score-only	Greedy	DDQN
C2	Event-shaped	Event (C2)	Greedy	DDQN
C3	Dueling+PER+SP	Event (C2)	Self-play	Dueling + PER
C4	Card Rush	Card Rush	Greedy	DDQN
C5	Noble Hunter	Noble Hunter	Greedy	DDQN
C6	Balanced Dense	Balanced Dense	Greedy	DDQN

All conditions: 300 000 episodes, evaluated every 1 000 episodes over 100 greedy-play games ($\varepsilon = 0$). Metrics: win rate, episode length, smoothed return, and card-purchase distribution. Full hyperparameter configuration is listed in Table 2.

Table 2: Training hyperparameters shared across all conditions (unless overridden by the condition configuration).

Hyperparameter	Value	Note
State dimension	213	Fixed-length normalized vector
Action dimension	48	Discrete, with legal mask
Hidden layer sizes	(256, 256, 128)	Shared trunk for standard DDQN
Optimizer	Adam	Learning rate 10^{-3}
LR step-down	$\div 10$	At episodes 200 k and 250 k
Discount factor γ	0.99	
Batch size	64	Uniform or PER sampling
Replay buffer capacity	10^5	Circular (FIFO)
Min buffer fill	5 000	Before any gradient step
Soft-update rate τ	0.005	Applied every gradient step
$\varepsilon_0, \varepsilon_{\min}, \lambda$	1.0, 0.05, 0.99997	Per-episode decay
Max steps per episode	200	Safety truncation
PER α, β_0	0.6, 0.4	C3 only; $\beta \rightarrow 1.0$ over 500 k steps
Self-play pool size	10 snapshots	C3 only; snapshot every 2000 ep
Gradient clip norm	1.0	All conditions
Win / loss terminal reward	± 10	All conditions

9 Contributions

- A Gym-compatible *Splendor* environment with legal action masking and a modular opponent interface (greedy, human, self-play).
- A masked DDQN implementation with optional Dueling architecture, PER, LayerNorm, AMP, and torch.compile—providing a reproducible baseline for sparse-reward board-game DRL.
- A self-play curriculum (C3) enabling progressive opponent difficulty without manual curriculum design.
- A real-time Pygame visualization enabling full game inspection in AI vs. AI and AI vs. Human modes.
- Empirical ablation across seven reward-shaping and architectural conditions, identifying both the cowardice pathology and the combination of techniques required to overcome it.

10 Faced Problems and Observations

During training and environment formulation, two critical challenges emerged that exposed fundamental issues in applying standard DRL to multi-player board games.

10.1 MDP Formulation and Action Space Masking

Accurate formulation required handling a dynamic action space whose legality depends on the current board state. Assigning a negative penalty for illegal moves caused the agent to gravitate toward inaction. Strict action masking—setting illegal logits to -10^4 before $\arg \max$ —resolved this completely.

However, masking interacts subtly with the discard rule. When a player holds more than 10 tokens, only discard actions are legal. The MDP treats each discard as an independent timestep returning zero reward. This creates a mandatory sub-phase within one game turn that the agent can intentionally trigger.

10.2 Discount-Exploitation Pathology

The most significant challenge was a classic RL exploit we term the *Cowardice Strategy*. With terminal loss penalty -10 and $\gamma = 0.99$, the discounted penalty at episode horizon H is -10×0.99^H . An agent that cannot reliably win finds it mathematically rational to maximize H :

$$-10 \times 0.99^{85} \approx -4.3 \quad > \quad -10 \times 0.99^{30} \approx -7.4$$

The agent discovered it could extend H by intentionally hoarding gems, forcing itself over the 10-token limit. Each mandatory discard action then counted as an extra timestep returning reward 0, converting one game turn into three RL steps and reducing the discounted loss. This pathology was most severe in early versions of C2 and C4, where intermediate rewards were insufficient to compete with the discounting incentive. All three fixes in Section 11 were required to eliminate it.

11 Implementations and Fixes

11.1 Advanced Architectural Enhancements

- **Dueling DQN:** Separates $V(s)$ and $A(s, a)$ estimation. The board state determines which turns are high-value independently of the specific action chosen.
- **Prioritized Experience Replay (PER):** Rare but critical transitions—card purchases, noble attractions, wins—receive higher replay frequency proportional to their TD error.
- **LayerNorm + Huber loss + gradient clipping:** Together stabilize training on the high-variance, long-horizon episodes characteristic of *Splendor*.

11.2 Self-Play Curriculum

Training C3 against past snapshots of itself eliminates the fixed-opponent exploit at the root: past versions of the agent also try to win, so extending episode length provides no benefit. Progressive difficulty emerges

naturally from the pool without manual curriculum design.

11.3 Reward Shaping Calibration

Four profiles were tested to address specific pathological incentives:

- **C4 (Card Rush)**: Zero gem reward; high card-buy rewards to incentivize immediate gem spending.
- **C5 (Noble Hunter)**: Double noble reward (+6.0) and elevated L3-card reward to target the fastest 15-point path.
- **C6 (Balanced Dense)**: All intermediate rewards amplified 2–5× for denser gradient signal.

12 Results and Analysis

We organize the results around the three problem statements from Section 2, then present the overall performance summary and ablation. All win-rate figures are evaluated against the fixed greedy opponent at $\varepsilon = 0$.

12.1 Problem Statement 1 — Sparse Rewards

Question: Does event-based reward shaping provide a denser learning signal that accelerates training beyond what score-only reward achieves?




Figure 3: Sample efficiency: episodes to first reach 50% win rate vs. the greedy opponent. C0 (trained vs. random) never reaches the threshold. All other conditions do; C4 (Card Rush, 144 k) and C3 (Dueling+PER+Self-play, 160 k) are most sample-efficient. Score-only C1 requires 298 k—the least efficient of all successful conditions, 154 k more than C4.

Figure 3 directly answers Problem Statement 1. Score-only C1 needs 298 000 episodes to reach 50% win rate—the largest budget of any successful condition. Every event-shaped variant (C2–C6) reaches the threshold faster:

- **C2 (event-shaped)**: 237 k — saving 61 k episodes over C1.
- **C6 (balanced dense)**: 208 k — all intermediate rewards amplified 2–5×.
- **C4 (card rush)**: 144 k — fastest overall; zero gem reward forces direct focus on card purchases.
- **C3 (Dueling+PER+self-play)**: 160 k — second fastest; combines event shaping with architectural improvements.

Answer — PS1 ✓: Event-based reward shaping is confirmed to densify the learning signal. All shaped conditions reach the 50% threshold more than 60 k episodes earlier than score-only training. The gain is largest when rewards are calibrated to the most strategically decisive actions (C4), confirming the Rinascimento hypothesis [1] in a model-free DRL setting.

12.2 Problem Statement 2 — Delayed Strategic Payoff

Question: Can the agent learn the multi-turn delayed strategy of building a gem engine first, then converting it to prestige?




Figure 4: Win-rate learning curves over 300 000 training episodes for all conditions. Each curve exhibits a clear S-shaped trajectory: a flat early phase (≈ 0 –10% win rate during initial exploration), a steep rise as the agent discovers the gem-collection-to-card-purchase conversion strategy, and a late-training plateau. This three-phase learning pattern directly mirrors the in-game delayed payoff structure of Splendor: the agent cannot accrue prestige without first investing multiple turns into gem accumulation. C4 (Card Rush) and C3 (Dueling+PER+Self-play) begin their steep rise earliest and reach 50% win rate at 144 k and 160 k episodes respectively; score-only C1 is the slowest at 298 k.

Figure 4 captures the learning-space analog of Splendor’s delayed payoff. The flat early phase (~ 0 –30 k episodes) corresponds to the agent randomly exploring actions without yet grasping that gem collection must precede card purchasing. The steep mid-phase rise occurs when the agent begins to internalize the multi-turn dependency: spending several zero-reward turns accumulating gems is a necessary prerequisite for the high-prestige card purchases that ultimately win the game. The plateau phase reflects policy convergence once this temporal structure is fully internalized.

Across all seven conditions, this S-shaped trajectory is consistent, confirming that the DDQN framework discovers the correct three-phase in-game structure — early-game gem engine building, mid-game card purchasing, late-game prestige completion — purely from experience, without any explicit hand-coded phase transition. Conditions with event-based rewards (C2–C6) enter the steep learning phase earlier than score-only C1, because intermediate card-purchase rewards provide gradient signal during the critical mid-game transition.

Answer — PS2 ✓: All trained agents successfully learn the delayed multi-turn strategy. The consistent S-shaped win-rate trajectories across all conditions are direct evidence that DDQN internalizes the long-horizon dependency: early gem collection yields no immediate reward but is an indispensable prerequisite for the prestige-gaining card purchases that win the game.

12.3 Problem Statement 3 — State-Dependent Legal Action Space

Question: Does legal action masking enable the agent to navigate the dynamically changing valid action set without attempting illegal moves?

The episode-length curves (Figure 5) provide behavioral evidence that legal action masking enables the agent to navigate the dynamic action space correctly. The mechanical guarantee is provided by the mask construction ($Q_{\text{masked}} = -10^4$ for all illegal actions), which ensures no gradient ever flows through an infeasible choice. Beyond this guarantee, the behavioral evidence is equally important:




Figure 5: Mean episode length (turns per game) over 300 000 training episodes for all conditions. After the initial random-policy phase, all conditions converge to stable game lengths of 25–35 turns—consistent with expert-level Splendor play. Critically, episode lengths do not grow unbounded in the final training range: this confirms that the discount-exploitation (cowardice) pathology described in Section 10 is fully resolved and that trained agents select strategically valid, legally coherent actions rather than deliberately inflating turn count via the discard sub-phase.

- **Stable game lengths (25–35 turns)** in all final policies confirm that agents are not exploiting the discard sub-phase to artificially inflate episode length. Early in training, some conditions show longer episodes (~ 45 – 60 turns) while the cowardice exploit is active; this collapses as reward shaping and the fixes in Section 11 take effect.
- **Convergence shape.** The steep drop in episode length during the mid-training phase (coinciding with the win-rate rise in Figure 4) indicates that the agent transitions from defensive gem-hoarding to decisive card-purchasing—the correct strategic response to the dynamic action set.
- **C4 converges fastest.** The Card Rush reward profile’s zero gem-collection weight eliminates any incentive to hoard gems, causing C4 to reach stable episode lengths earliest—consistent with its fastest approach to 50% win rate.

Answer — PS3 ✓: Legal action masking completely solves the feasibility problem. Beyond mechanical correctness, the convergence of episode lengths to the 25–35 turn range demonstrates that agents learn to operate *strategically* within the changing action set, selecting gem-take, buy, and reserve actions in proportions consistent with effective Splendor play.

12.4 Overall Performance Summary

Figure 6 summarizes peak performance. **Every condition trained against the greedy opponent (C1–C6) reaches above 50% peak win rate**, confirming that the full DRL pipeline—environment, legal masking, replay, and reward shaping—successfully produces agents that beat the greedy heuristic. The cowardice pathology described in Section 10 was fully mitigated by the fixes in Section 11, and episode lengths in the final trained policies remain stable near 25–30 steps.

A key insight from Table 3 is the distinction between *peak* win rate and *final-phase stability*. C1 (score-only) achieves a 52% peak but only 21.0% stable win rate in the last 50 k episodes and C0 stabilizes at just 5.5%—both trained policies show high variance or late-training degradation. In contrast, C4 achieves the highest final-phase stability ($39.9\% \pm 9.3\%$), confirming that Card Rush reward calibration produces a *consistently reliable* policy. C3 (Dueling+PER+self-play) reaches 32.3% stability—lower than C4 when measured against the fixed greedy opponent, though self-play introduces continued improvement in cross-condition generalization (Section 12.5).




Figure 6: Peak performance comparison: best win rate vs. greedy across all 300 000 training episodes. All conditions except C0 (trained vs. random) surpass the 50% competency threshold. C3 (Dueling+PER+Self-play) and C4 (Card Rush) jointly achieve the highest peak win rate of 60%.

Table 3: Final performance summary. Final-phase stability = mean $\pm$ std win rate over the last 50 k training episodes.

Cond.	Description	Peak Win Rate	Ep. to 50%	Final Stability
C0	Score-only vs. Random	39%	N/A	$5.5 \pm 4.3\%$
C1	Score-only vs. Greedy	52%	298 k	$21.0 \pm 11.9\%$
C2	Event-shaped vs. Greedy	59%	237 k	$36.6 \pm 11.4\%$
C3	Dueling+PER+SP	60%	160 k	$32.3 \pm 8.8\%$
C4	Card Rush vs. Greedy	60%	144 k	$39.9 \pm 9.3\%$
C5	Noble Hunter vs. Greedy	57%	263 k	$33.8 \pm 11.3\%$
C6	Balanced Dense vs. Greedy	56%	208 k	$38.9 \pm 8.9\%$

12.5 Cross-Condition Tournament




Figure 7: Round-robin cross-condition tournament. *Left:* win-rate matrix (row = player, column = opponent); each entry is the win rate of the row agent vs. the column agent over 100 games. *Right:* ranking by average win rate vs. all other conditions. C4 (Card Rush) ranks #1 with 61.1% average win rate, ahead of C3 (58.4%) and C6 (57.9%). C0, trained only against a random opponent, ranks last at 38.8%.

The tournament (Figure 7) evaluates generalization by pitting every trained agent directly against every other. C4 (Card Rush) emerges as the strongest overall condition with a 61.1% average win rate, ranking #1 across all head-to-head matchups. C3 (Dueling+PER+Self-play) places #2 at 58.4%—the self-play curriculum gives C3 strong cross-condition robustness even against agents it has never seen during training. Notably, the

gap between C2 through C6 is narrow (57.2–61.1%), while C0 is clearly separated at 38.8%, confirming that training against only a random opponent produces a non-generalizing policy.




Figure 8: Best-checkpoint vs. final-checkpoint win rate, and late-training degradation. *Left:* best checkpoint (dark bar) vs. final checkpoint (hatched bar) for each condition. *Right:* degradation magnitude—C0 (+31.5%) and C2 (+35.0%) show severe regression; C1 and C6 are stable ($\leq 3\%$); C3 uniquely *improves* (-7.5%) as self-play progressively strengthens the opponent. C4 shows moderate regression (+20.0%) but retains the highest best-checkpoint win rate (55.5%).

Figure 8 adds a longitudinal perspective: C4 achieves the highest single-checkpoint peak (55.5%) but regresses 20,pp by the final episode. C3 is the only condition that *continues to improve* throughout training (final 49.0% > best-window 41.5%), driven by the self-play curriculum progressively raising opponent difficulty. C6 uniquely sustains its peak with zero degradation.

The ablation ladder (Figure ??) isolates the marginal contribution of each design choice:

Opponent type — C0 vs. C1 (+13,pp peak). Training against the greedy rather than random opponent raises peak win rate from 39% to 52%. A harder training distribution produces a stronger policy.

Reward shaping — C1 vs. C2 (+7,pp peak, -61 ,k episodes, +15.6,pp stability). Event shaping provides the largest single gain in both speed and consistency. The stability jump from 21.0% to 36.6% confirms that score-only training is highly unstable in late training.

Reward calibration — C2 vs. C4, C5, C6 (± 1 –3,pp, -30 to -93 ,k episodes). Different weight profiles shift sample efficiency substantially while keeping peak win rate within 3,pp. Calibration controls *how fast* the agent learns, not *how good* it gets.

Architecture + curriculum — C2 vs. C3 (+1,pp peak, -77 ,k episodes). Dueling DQN, PER, and self-play improve speed but yield lower fixed-opponent stability (32.3% vs. C2’s 36.6%). This tradeoff pays off in the tournament (Section 12.5), where C3 places 2 overall.

Overall winner — C4 across three metrics. C4 reaches 50% win rate fastest (144,k), holds the highest final-phase stability (39.9%), and ranks 1 in the cross-condition tournament (61.1%). Its only drawback is moderate late-training regression, while C3 keeps improving and C6 holds steady.

13 Limitations

Fixed evaluation opponent. All conditions are evaluated exclusively against the same greedy opponent. A condition that beats the greedy heuristic may not generalize to human players or stronger algorithmic opponents. Mode 2 (AI vs. Human) provides qualitative validation but is not systematically logged.

Incomplete ablation matrix. The seven conditions do not form a fully factorial design. In particular, there is no condition combining self-play with a standard DDQN (without Dueling/PER), nor one using

Dueling+PER without self-play. The relative contribution of each C3 component therefore rests partly on theoretical argument rather than direct measurement.

No search-based baseline. The greedy opponent is a weak heuristic. An MCTS or minimax baseline would better characterize how close the 60% win rate is to the performance ceiling against a near-optimal policy.

Two-player only. The standard *Splendor* game supports 2–4 players. The current MDP formulation assumes exactly two players; extending to multi-player requires reformulating the turn structure and the opponent interface.

Action space approximation. The gem-discard sub-phase is modeled as individual single-color discards rather than simultaneous multi-color selection, artificially inflating the number of RL timesteps per game turn. This is the structural asymmetry exploited by early cowardice pathology.

14 Conclusion

This project investigated DRL for *Splendor* across seven conditions, showing that event-based reward shaping, Dueling DQN, PER, and self-play curriculum together address all three original problem statements.

1. **All three problem statements resolved.** Action masking eliminates illegal moves (PS3); shaped rewards accelerate learning by up to 154k episodes over score-only (PS1); all agents learn the gem-engine-then-prestige strategy (PS2).
2. **Every greedy-trained agent exceeds 50% win rate.** Peak rates range from 52% (C1) to 60% (C3, C4), confirming the DRL pipeline reliably beats the greedy heuristic across all reward designs.
3. **C4 (Card Rush) is the strongest practical condition.** It reaches 50% fastest (144k episodes), holds the highest final-phase stability ($39.9\% \pm 9.3\%$), and ranks #1 in the cross-condition tournament (61.1%). C3 ranks #2 (58.4%) and uniquely keeps improving throughout training via its self-play curriculum.

The modular, Gym-compatible system provides a reproducible platform for future work—stronger opponents, multi-player extensions, or search-based policy improvement.

References

- [1] I. Bravi and S. M. Lucas, “Rinascimento: using event-value functions for playing Splendor,” in *Proc. IEEE Conf. on Games (CoG)*, 2020.
- [2] V. Mnih *et al.*, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
- [3] H. van Hasselt, A. Guez, and D. Silver, “Deep reinforcement learning with Double Q-learning,” in *Proc. AAAI Conf. on Artificial Intelligence*, 2016.
- [4] Z. Wang *et al.*, “Dueling network architectures for deep reinforcement learning,” in *Proc. ICML*, 2016.
- [5] T. Schaul, J. Quan, I. Antonoglou, and D. Silver, “Prioritized experience replay,” in *Proc. ICLR*, 2016.
- [6] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. MIT Press, 2018.
- [7] S. Huang and S. Ontañón, “A closer look at invalid action masking in policy gradient algorithms,” in *Proc. FLAIRS*, vol. 35, 2022.